

Gaussian Filtering and Edges

Pablo Mesejo

pmesejo@go.ugr.es

Universidad de Granada

Departamento de Ciencias de la Computación e Inteligencia Artificial




UNIVERSIDAD
DE GRANADA





DaSCI

Instituto Andaluz de Investigación en
Data Science and Computational Intelligence

Smoothing with box filter



Gaussian kernel






An isotropic (circularly symmetric) filter


$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

- The coefficients are a 2D Gaussian.
- Gives more weight at the central pixels and less weight to the neighbors.
- The farther away the neighbors, the smaller the weight.


Gaussian filters



$$\sigma = 1$$



$$\sigma = 5$$




$$\sigma = 10$$



$$\sigma = 30$$

Smoothing with a Gaussian

Parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



```
for sigma in np.arange(1, 10, 3):  
    kernel=cv2.getGaussianKernel(ksize, sigma)  
    outim=cv2.filter2D(im,-1,kernel) #correlation  
    cv2.imshow(outim)
```


If you use Google Colab, from
`google.colab.patches import`
`cv2_imshow`

When `ddepth=-1`, the output image will
have the same depth as the source

Mean vs. Gaussian filtering



**The box filter provides
edgy artifacts!**




**Gaussian is a better low-pass
filter → it does not
create high-frequency
artifacts.**




Gaussian filter


- What parameters matter here?
 - **Size** of kernel or mask
 - **Variance** of Gaussian determines extent of smoothing



$\sigma = 5$ with
10 x 10
kernel



$\sigma = 5$ with
30 x 30
kernel

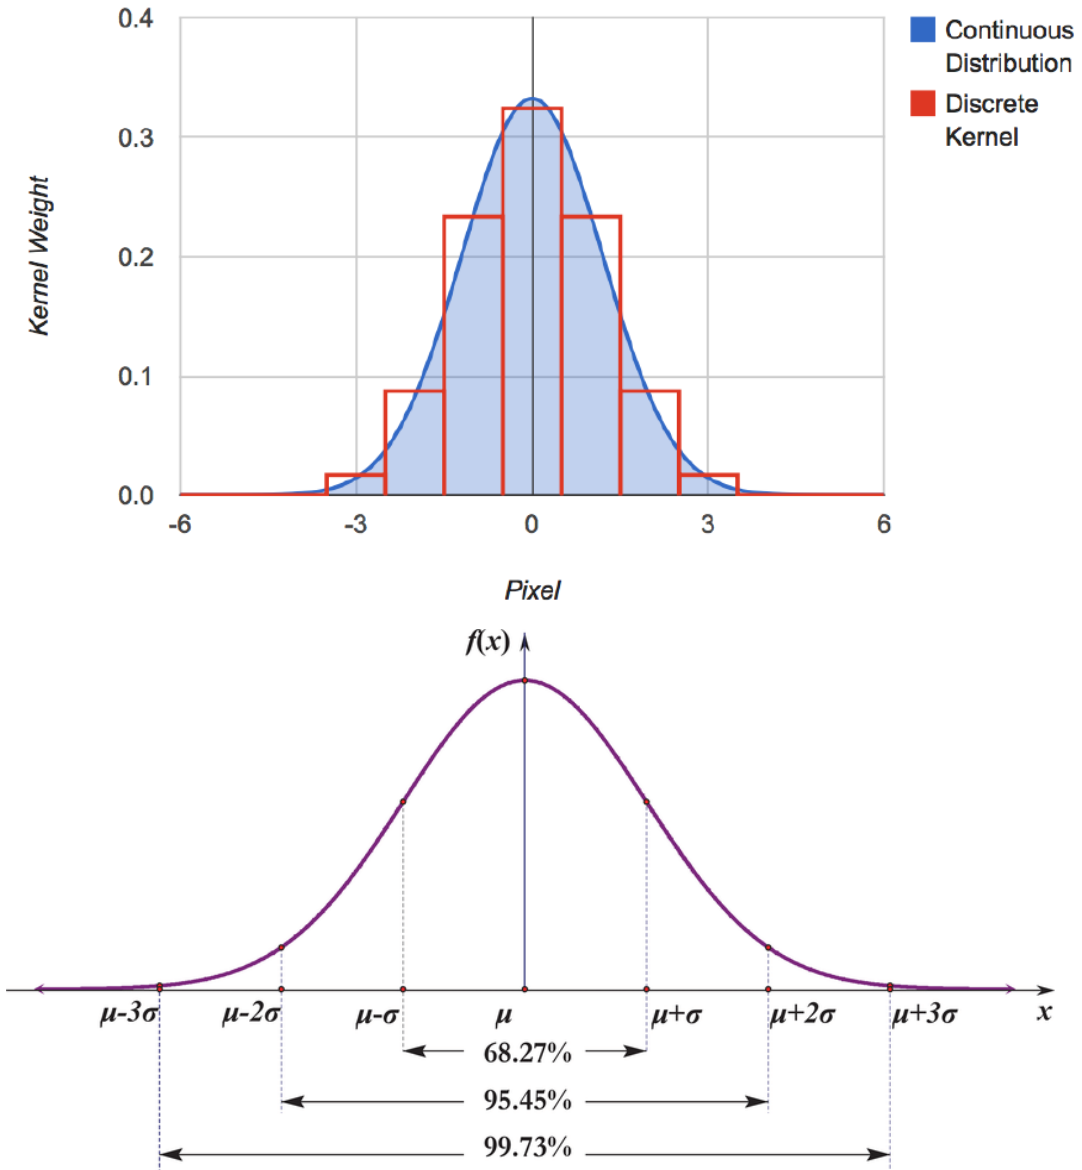


$\sigma = 2$ with
30 x 30
kernel

Gaussian discretization

We want to discretize a Gaussian.

We know that **almost all values are within 3 standard deviations of the mean** → It makes sense to use, by default, 3σ .



Gaussian discretization

$$f(x, \sigma) = c \cdot e^{-\frac{x^2}{2\sigma^2}}$$

We ignore this normalizing constant ($1/\sigma\sqrt{2\pi}$)! The shape of the filter is the same! We later normalize the kernel, making the sum of the coefficients equal to 1

Gaussian discrete mask 1D:

$$[f(-k), f(-k+1), \dots, f(0), \dots, f(k-1), f(k)], \quad k \text{ an integer}$$

According to the Gaussian properties, $\min(k) \geq 3\sigma$ provides samples covering more than 99% of the area under the curve

**If $\sigma = 1 \rightarrow$ kernel size is 7
(ie, $k = 3$)**

$$\text{We enforce } \sum_{i=-k}^k f(i) = 1$$

The mask must add up to 1 \rightarrow you must divide your mask by the sum of its coefficients

Why?? 🤔

Gaussian discretization

We can go from σ to mask size (T), or viceversa, in order to get the Gaussian mask

$$\text{If } T \text{ is the mask size, } T = 2 \cdot k + 1 \rightarrow k = (T-1)/2 \rightarrow (T-1)/2 = 3 \cdot \sigma \rightarrow 2 \cdot [3 \cdot \sigma] + 1 = T$$

Let's see an example. We have $T=5$:

- We'd have 5 positions: $[f(-2), f(-1), f(0), f(1), f(2)]$
- We compute the suitable sigma: $(T-1)/2 = 3 \cdot \sigma \rightarrow (5-1)/2 = 3 \cdot \sigma \rightarrow 2 = 3 \cdot \sigma \rightarrow \sigma = 2/3 = 0.67$
- We compute the values of the mask substituting $x = \{-2, -1, 0, 1, 2\}$ on $f(x) = e^{-\frac{x^2}{2 \cdot 0.67^2}}$

Once we have σ and k we can discretize the mask by applying the Gaussian function (or its derivatives)

Gaussian discretization: finite approximation

Convolving a box filter with itself yields an approximation to a Gaussian kernel (easy way to construct a Gaussian kernel):

$$\frac{1}{2} \begin{bmatrix} 1 & 1 \end{bmatrix} * \frac{1}{2} \begin{bmatrix} 1 & 1 \end{bmatrix} = \frac{1}{4} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$
$$\frac{1}{4} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \frac{1}{4} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} = \frac{1}{16} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

These are the odd rows of the binomial (Pascal's) triangle:

$k=1, \sigma^2=0.5, \sigma=0.7$

$k=2, \sigma^2=1, \sigma=1$

		1					
		1		1			
		1	2	1			
	1		3		3		1
	1	4	6	4	1		

(2k+1)th row approximates
Gaussian with $\sigma^2=k/2$

Check Stan Birchfield's notes for more
information regarding the binomial and
trinomial triangle to generate Gaussian
approximations


Gaussian Properties

- Remove “high-frequency” components from the image (*low-pass filter*). All values are positive and sum up to 1.
 - Images become smoother.
- Completely described by 1st (mean) and 2nd (variance) order statistics (moments).
- The Fourier Transform of a Gaussian function is also a Gaussian

$$f(x) = e^{-\frac{x^2}{2\sigma^2}} \longrightarrow \mathcal{F}[f(x)] = \sqrt{2\pi}\sigma e^{-2\pi^2\sigma^2\omega^2}$$

Gaussian Properties

- Closed under convolution (convolution of two Gaussians is another Gaussian)



- When you convolve a Gaussian function (G_1) with another Gaussian function (G_2), the result is still a Gaussian function (G)
- Convolving twice with Gaussian kernel of width σ is the same as convolving once with kernel of width $\sqrt{\sigma^2 + \sigma^2} = \sqrt{2\sigma^2} = \sigma\sqrt{2}$
- In general, n repeated convolutions with σ Gaussian approximates single convolution with $\sigma\sqrt{n}$ Gaussian

$$G(x, y|\sigma) = G\left(x, y \mid \sqrt{\sigma_1^2 + \sigma_2^2}\right) = G_1(x, y|\sigma_1) * G_2(x, y|\sigma_2)$$

Gaussian Properties

- Separable (efficient)
 - Factors into product of two 1D Gaussians
 - Example:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$



```
import scipy.signal as sp
import numpy as np

np.outer(np.array([1,2,1]), np.array([1,2,1]))

sp.convolve2d(np.array([[0,1,0],[0,2,0],[0,1,0]]),
              np.array([[0,0,0],[1,2,1],[0,0,0]]), mode='same')
```

Outer product:

$$\mathbf{u} = \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_m \end{bmatrix}, \quad \mathbf{v} = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} \quad \mathbf{u} \otimes \mathbf{v} = \mathbf{A} = \begin{bmatrix} u_1 v_1 & u_1 v_2 & \dots & u_1 v_n \\ u_2 v_1 & u_2 v_2 & \dots & u_2 v_n \\ \vdots & \vdots & \ddots & \vdots \\ u_m v_1 & u_m v_2 & \dots & u_m v_n \end{bmatrix}$$

Separability of the Gaussian

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability example

**2D filtering
(center location only)**

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} = 2 + 6 + 3 = 11 \\ = 6 + 20 + 10 = 36 \\ = 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

**The filter factors
into a product of 1D
filters:**

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

**Perform filtering
along rows:**

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & 11 & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array}$$



**Followed by filtering
along the remaining column:**

$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline & 11 & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & 65 & \\ \hline & & \\ \hline \end{array}$$

Separability example

We combine information from different rows
(kernel Y - Coefficients for filtering each column)

We combine information from different columns
(kernel X - Coefficients for filtering each row)



Input Image

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 10 & 20 & 30 & 0 \\ 0 & 40 & 50 & 60 & 0 \\ 0 & 70 & 80 & 90 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Kernel

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

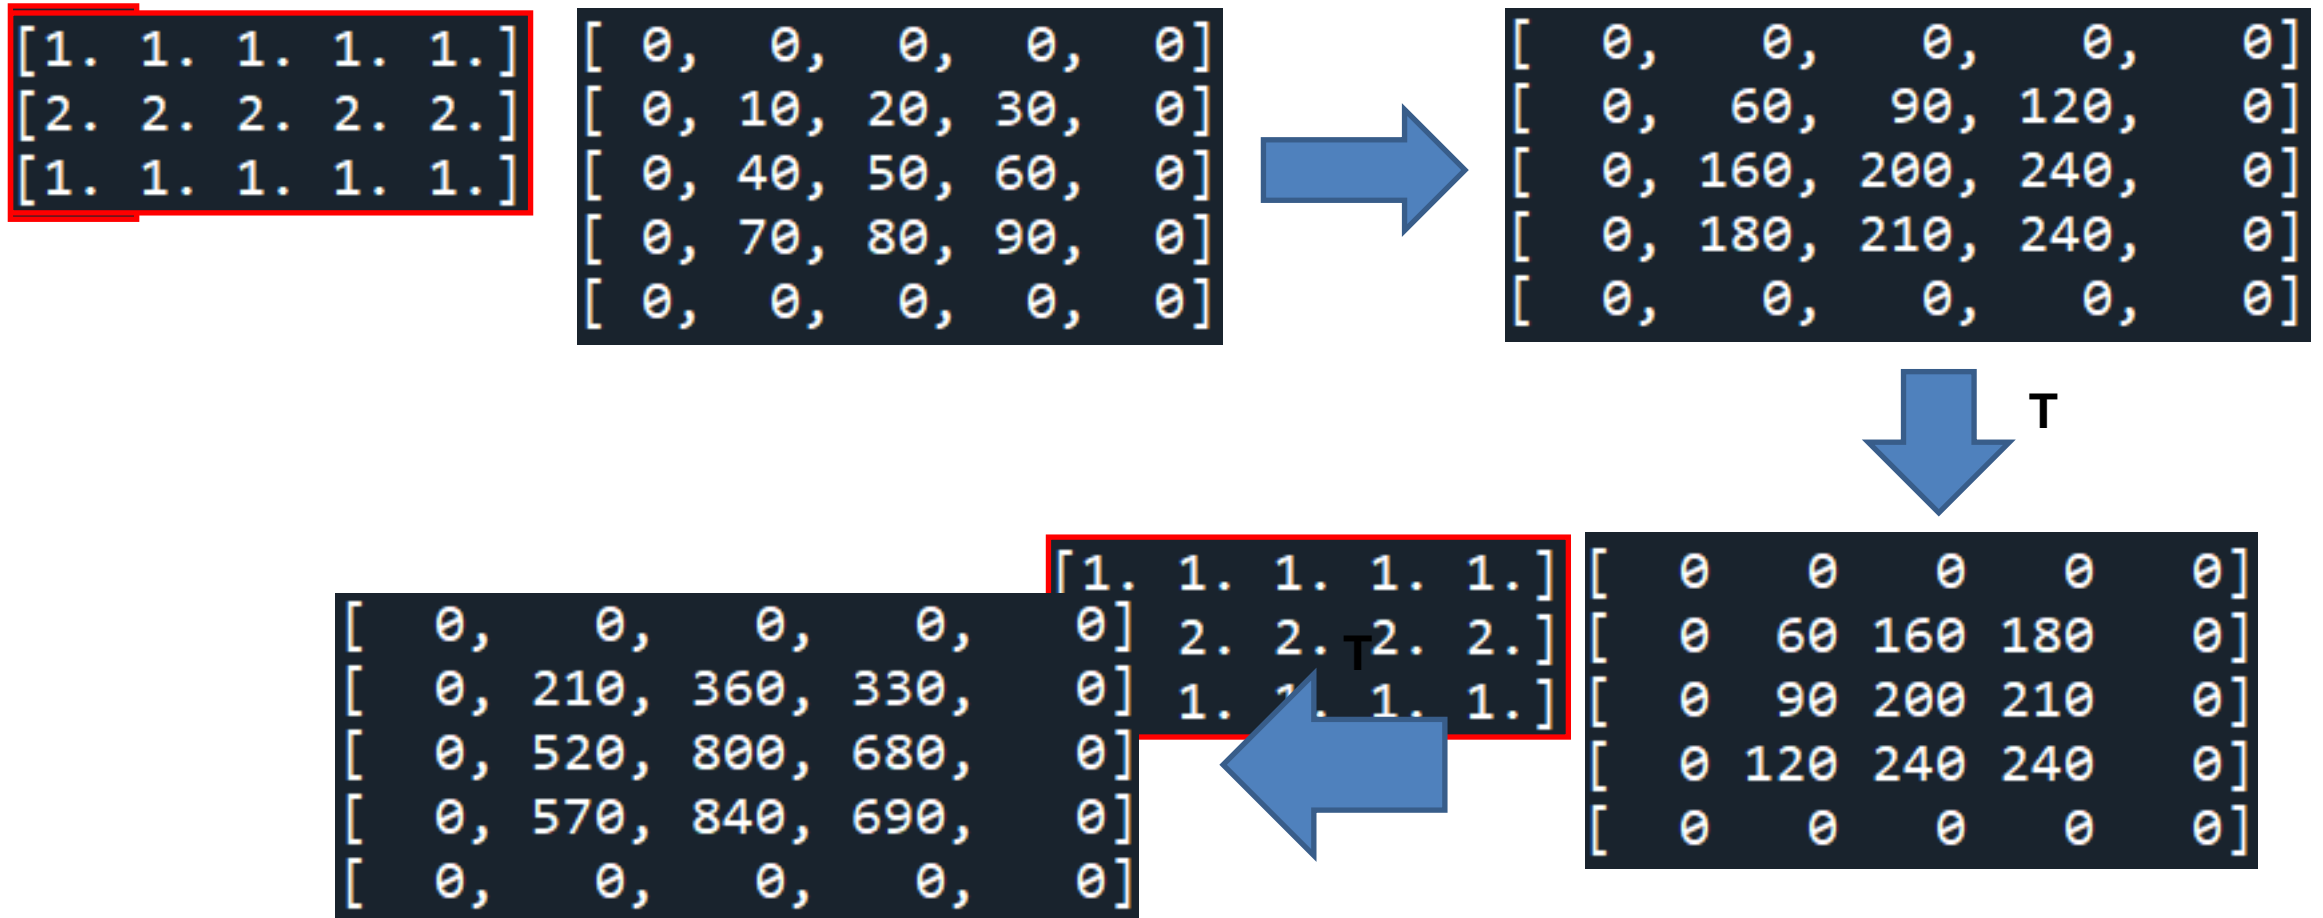
Separable Kernel

http://www.songho.ca/dsp/convolution/convolution2d_separable.html



$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 210 & 360 & 330 & 0 \\ 0 & 520 & 800 & 680 & 0 \\ 0 & 570 & 840 & 690 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Separability example (technical note)

It is not necessary to always go through all the rows and columns (4 for-loops)!!!
For each 1D convolution, we only need to traverse the rows once!!!



Why is separability useful in practice?

- Separability means that a 2D convolution can be reduced to two 1D convolutions (one among rows and one among columns)
- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?
 - $O(n^2 m^2)$  We move the $m \times m$ filter through the whole image. So, we do $m \cdot m$ operations for each of the $n \cdot n$ pixels in the image: $m^2 \cdot n^2$ operations
- What if the kernel is separable?
 - $O(n^2 m)$  We move the $1 \times m$ filter through the whole image. So, we do m operations for each of the $n \cdot n$ pixels in the image (one per rows and another per columns): $m \cdot n \cdot n + m \cdot n \cdot n = 2 \cdot (m \cdot n^2)$

Example: 512x512 image, 11x11 filter

- 1 convolution 2D: $512 \cdot 512 \cdot 11 \cdot 11$ ~32M multiplications
- 2 convolutions 1D: $512 \cdot 512 \cdot 11 + 512 \cdot 512 \cdot 11$ ~6M multiplications

More ideas about separability

- Allowed because convolution is associative ($f * (g_1 * g_2) = (f * g_1) * g_2$)
- Some 2D kernels can be directly decomposed into 1D kernels:

$$\begin{aligned} \frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} &= \frac{1}{4} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \frac{1}{4} [1 \quad 2 \quad 1] \\ &= \frac{1}{4} [1 \quad 2 \quad 1] * \frac{1}{4} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \end{aligned}$$

$$\frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} * \frac{1}{3} [1 \quad 1 \quad 1] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- A 2D kernel is *separable* iff all rows / columns are linearly dependent

– This filter $\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$ is not separable:

$$\alpha \cdot [0, 1, 0] + \beta \cdot [1, -4, 1] = [0, 0, 0] \iff \alpha = 0 \wedge \beta = 0$$

→ Linearly independent vectors!!

Note: any 2D-mask admits a SVD decomposition (*low-rank approximations*): sum of several separable kernels.

<https://bartwronski.com/2020/02/03/separate-your-filters-svd-and-low-rank-approximation-of-image-filters/>


Computing derivatives:
high-frequency features
(high-pass filters)

Edges

	24	25	26	27	28	29	30	31	32	33	34	35
24	15	20	23	15	22	16	14	17	13	20	14	14
25	21	13	14	15	18	14	17	23	21	17	67	61
26	25	12	16	19	22	22	20	14	25	64	69	63
27	24	19	15	26	27	25	15	55	59	52	61	58
28	13	21	27	20	18	28	54	62	66	56	53	65
29	13	14	24	20	23	68	56	69	61	57	69	62
30	14	28	18	28	63	67	68	57	67	67	69	66
31	24	18	22	59	55	66	59	55	68	54	56	56
32	26	23	23	58	55	57	64	63	57	58	54	56
33	19	22	68	61	56	64	69	56	55	68	53	57
34	22	55	62	57	62	59	59	64	55	68	67	57
35	19	54	60	60	57	55	61	56	57	61	55	60




Edges





Edges and High Frequencies

- Frequency in 2D images:
 - rate of change of gray levels in spatial terms
 - If it “involves” many pixels to make a change
→ low frequency
 - If it “involves” few pixels to make a change (i.e., the change is abrupt)
→ high frequency
- We’ll come back to this later!



Edge detection



- **Convert a 2D image into a set of curves**
 - Extracts salient features of the scene
 - More compact than pixels

Importance of intensity edges



from Walther et al., Simple line drawings suffice for functional MRI decoding of natural scene categories, PNAS 2011

→ *Much information is retained by the edges*


Importance of intensity edges



from Walther et al., Simple line drawings suffice for functional MRI decoding of natural scene categories, PNAS 2011


→ *Much information is retained by the edges*

Origin of edges



- Edges are caused by a variety of factors

Images as functions...

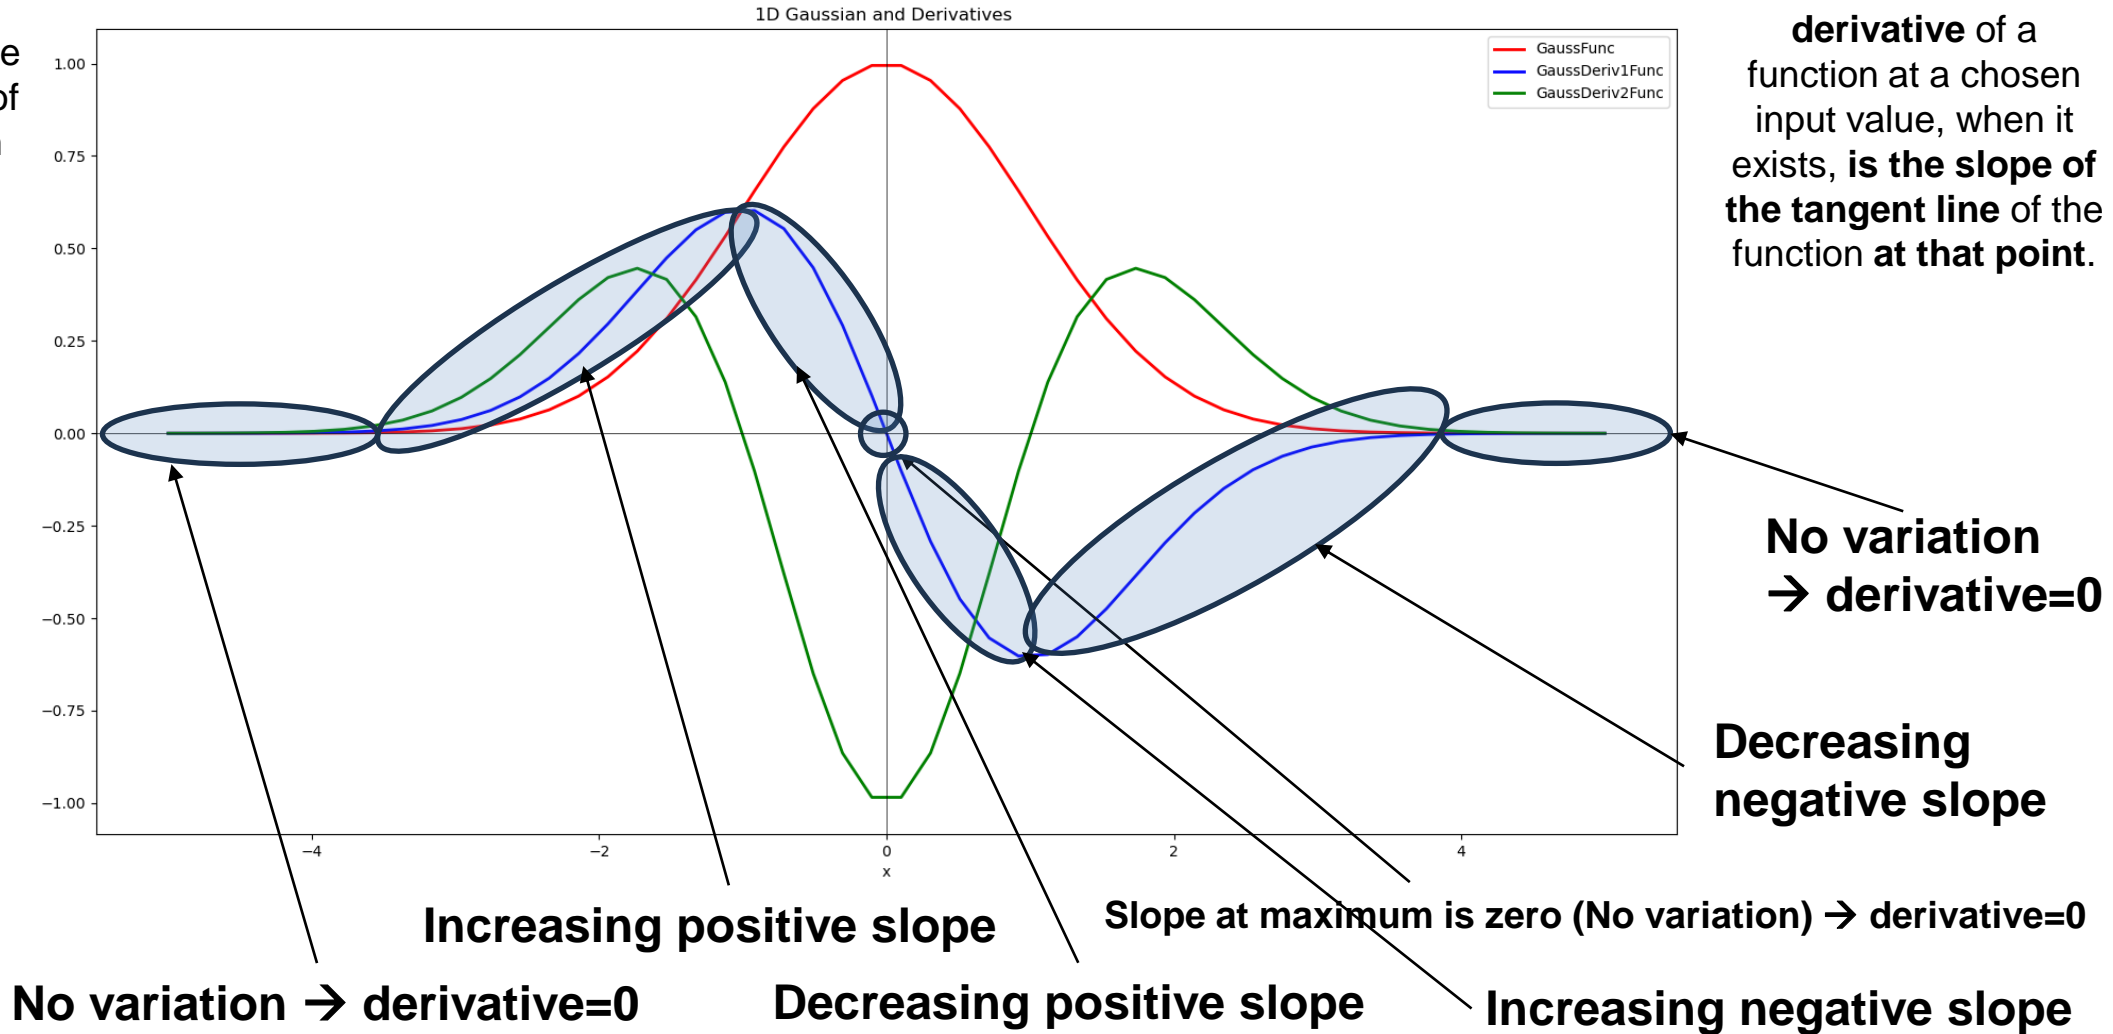


- Edges look like steep cliffs

Remembering the concept of derivative

The derivative shows the **sensitivity of change** of a function's output with respect to the input

Remember that the **derivative** of a function at a chosen input value, when it exists, **is the slope of the tangent line of the function at that point.**



Let's pay attention to the **blue** line (first derivative of the **red** one).
This will help us to better understand the next slides.

Characterizing edges

- An edge is a place of *rapid change* in the image intensity function




Image derivatives

For a 2D function, $f(x, y)$, the partial derivative is

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate the derivative using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

We compute derivatives through the **difference between intensities of adjacent pixels**.

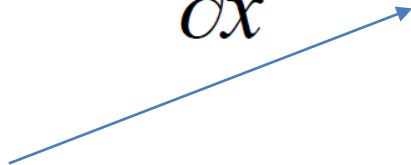


Image derivatives

- **How can we differentiate a digital image $F[x,y]$?**
 - **Take discrete derivative (finite difference)**

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$

How would you implement this as a linear filter?

$$\frac{\partial f}{\partial x} : \begin{array}{|c|c|c|} \hline & & \\ \hline 1 & -1 & \\ \hline & & \\ \hline \end{array}$$

H_x

$$\frac{\partial f}{\partial y} : \begin{array}{|c|c|c|} \hline & & \\ \hline & -1 & \\ \hline & 1 & \\ \hline \end{array}$$

H_y

Image derivatives

- Odd length avoids undesirable shifting of signal

$$\begin{bmatrix} 1 & \underline{-1} \end{bmatrix}$$

Forward difference

$$\begin{bmatrix} 4 & 7 & 8 & 1 & 3 \end{bmatrix} * \begin{bmatrix} 1 & \underline{-1} \end{bmatrix} = \begin{bmatrix} 3 & 1 & -7 & 2 & 0 \end{bmatrix}$$

(Don't forget to flip kernel before convolving)

The position to be replaced appears underlined.

In this example, we reflect the boundary. So, the value outside the image would be 3: $-1*3+1*3=0$

$$\begin{bmatrix} \underline{1} & -1 \end{bmatrix}$$

Backward difference

$$\begin{bmatrix} 4 & 7 & 8 & 1 & 3 \end{bmatrix} * \begin{bmatrix} \underline{1} & -1 \end{bmatrix} = \begin{bmatrix} 0 & 3 & 1 & -7 & 2 \end{bmatrix}$$

- To avoid undesirable shift, use central difference kernel

$$\frac{1}{2} \begin{bmatrix} 1 & \underline{0} & -1 \end{bmatrix}$$

Central difference

Example:

$$\begin{bmatrix} 4 & 7 & 8 & 1 & 3 \end{bmatrix} * \frac{1}{2} \begin{bmatrix} 1 & \underline{0} & -1 \end{bmatrix} = \frac{1}{2} \begin{bmatrix} 3 & 4 & -6 & -5 & 2 \end{bmatrix}$$

We don't really care about this factor.
We just want to detect differences!

Forward	$f'(x) \approx \frac{f(x+h) - f(x)}{h}$
Backward	$f'(x) \approx \frac{f(x) - f(x-h)}{h}$
Central	$f'(x) \approx \frac{f(x+0.5h) - f(x-0.5h)}{h}$

$h=2$

Image derivatives

- Another perspective to get central differences:

$$f_F'(x) \approx \frac{1}{h}(f(x+h) - f(x)) \text{ (forward)} \quad f_B'(x) \approx \frac{1}{h}(f(x) - f(x-h)) \text{ (backward)}$$

- First order central (average):

$$f_c'(x) \approx \frac{1}{2}(f_F' + f_B') = \frac{1}{2h}(f(x+h) - f(x-h)) \longrightarrow \frac{1}{2} \begin{array}{|c|c|c|} \hline \text{mask: } h=1 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

- Second order central:




$$f_c''(x) \approx \frac{f'(x+0.5h) - f'(x-0.5h)}{h} = \frac{\frac{f(x+h) - f(x)}{h} - \frac{f(x) - f(x-h)}{h}}{h} \longrightarrow \begin{array}{|c|c|c|} \hline \text{mask: } h=1 \\ \hline 1 & -2 & 1 \\ \hline \end{array}$$
$$= \frac{1}{h^2}(f(x+h) - 2f(x) + f(x-h))$$

In derivative masks, weight normalization is to 0

Image gradient

- The *gradient* of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The gradient points in the direction of most rapid increase in intensity




The *edge strength* is given by the **gradient magnitude**:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$


The **gradient direction** is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$



Decomposing the Sobel filter

Image gradient



Horizontal derivative






$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

Vertical derivative

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix}$$

Gradient magnitude

$$G = \sqrt{G_x^2 + G_y^2}$$



Visual examples taken from <https://stackoverflow.com/questions/19815732/what-is-the-gradient-orientation-and-gradient-magnitude> and https://en.wikipedia.org/wiki/Sobel_operator

Note: These filters give negative values, because derivatives can be negative. Then, at coding level, you have to “transfer” those values to [0,255] or [0,1].

Image gradient

Is there any difference if we
apply cross-correlation or
convolution with this kernel??



Horizontal derivative

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

Vertical derivative

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix}$$

Gradient magnitude

$$G = \sqrt{G_x^2 + G_y^2}$$

Gradient orientation

$$\Theta = \text{atan2}(G_y, G_x)$$







Image gradient




Which one represents the gradient magnitude?

Which one shows changes with respect to x?

Which one shows changes with respect to y?




Effects of noise




Noisy input image

$$f(x)$$



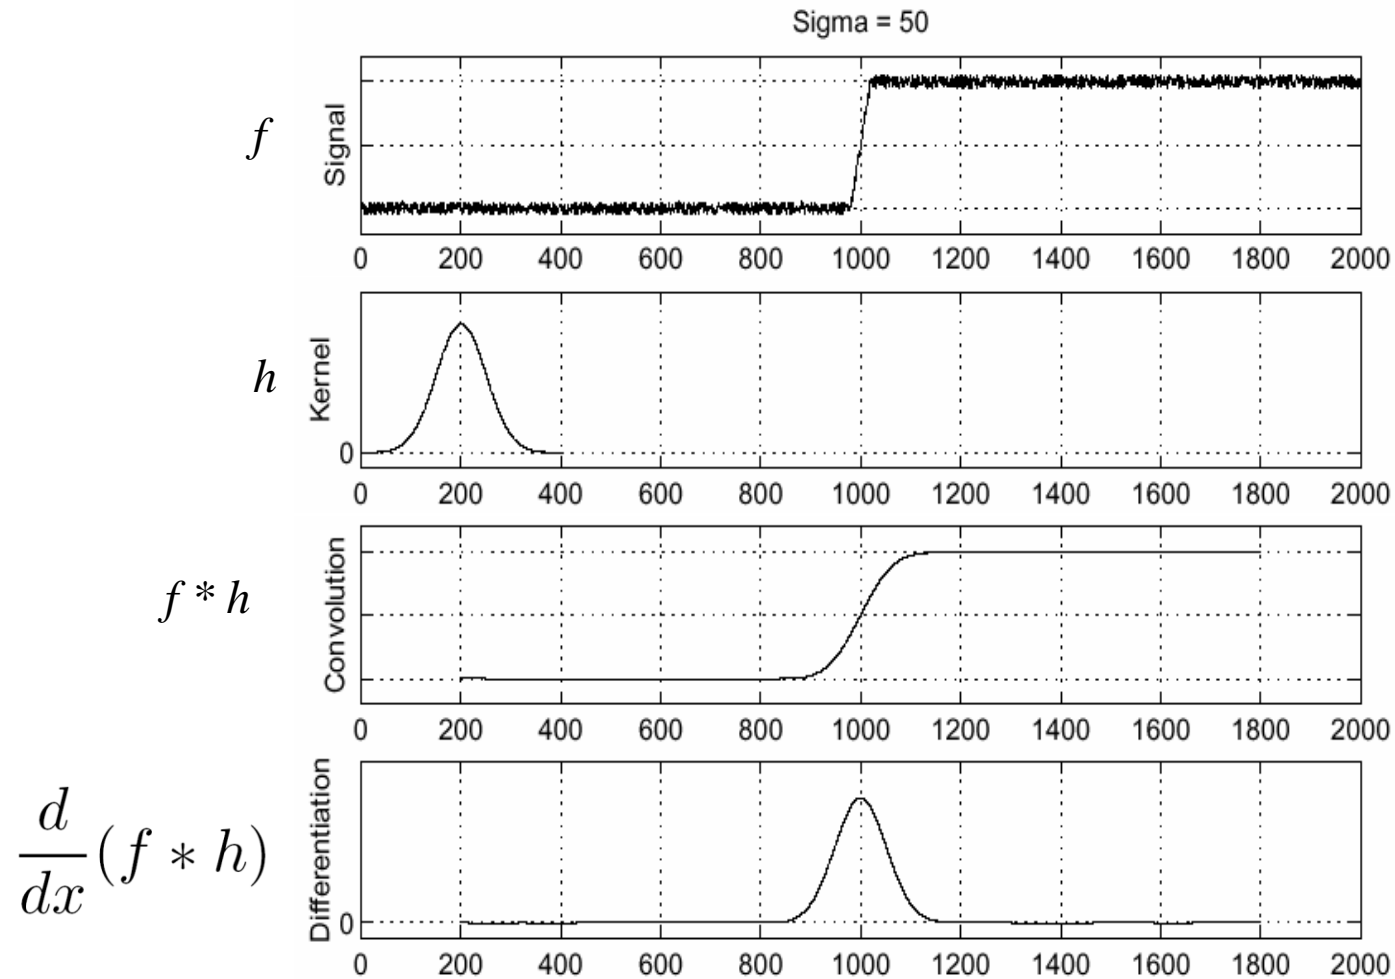
$$\frac{d}{dx}f(x)$$



**Derivatives
amplify noise**
(high-frequency
components)!

Where is the edge?

Solution: smooth first



How many convolutions do we have here??



To find edges, look for peaks $\frac{d}{dx}(f * h)$

Source: S. Seitz

Associative property of convolution

- Differentiation is convolution, and convolution is associative:

$$\frac{d}{dx}(f * h) = f * \frac{d}{dx}h$$

- This saves us one operation:

Computing derivatives implies smoothing (in one direction) and differentiating (in the other)


Examples:

$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} -1 & 0 & 1 \end{bmatrix} \Rightarrow \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$


1D Gaussian filter x-derivative Sobel - x

$$\frac{1}{2} \begin{bmatrix} 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & -1 \end{bmatrix} = \frac{1}{2} \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$


f



$\frac{d}{dx}h$



$f * \frac{d}{dx}h$




The 1D Gaussian and its derivatives

$$g_{\sigma}(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right)$$

$$g'_{\sigma}(x) = \frac{\partial g_{\sigma}(x)}{\partial x} = -\frac{x}{\sigma^2} \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) = -\frac{x}{\sigma^2} g_{\sigma}(x)$$

$$g''_{\sigma}(x) = \frac{\partial^2 g_{\sigma}(x)}{\partial x^2} = \left(\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2}\right) \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) = \left(\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2}\right) g_{\sigma}(x)$$




The 1D Gaussian and its derivatives




Intensity profile of an input image



After noise removal




First-order derivative



Second-order derivative


Main advantage of second derivatives? They allow you to locate the edge more precisely!

2D edge detection filters



Gaussian


$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$





derivative of Gaussian (x)

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$


Derivative of Gaussian filter



x-direction




y-direction



Derivative of Gaussian filter

Is the sign really important?? 🤔



Common 2D Gauss derivative kernels

Prewitt

$$Prewitt_x = \frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} * \frac{1}{2} \begin{bmatrix} 1 & 0 & -1 \end{bmatrix} = \frac{1}{6} \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

$$Prewitt_y = \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \end{bmatrix} * \frac{1}{2} \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} = \frac{1}{6} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

Sobel

$$Sobel_x = \frac{1}{4} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \frac{1}{2} \begin{bmatrix} 1 & 0 & -1 \end{bmatrix} = \frac{1}{8} \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$Sobel_y = \frac{1}{4} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \frac{1}{2} \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} = \frac{1}{8} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

The standard definition of the Sobel operator omits the 1/8 term

– doesn't make a difference for edge detection

Scharr

$$Scharr_x = \frac{1}{16} \begin{bmatrix} 3 \\ 10 \\ 3 \end{bmatrix} * \frac{1}{2} \begin{bmatrix} 1 & 0 & -1 \end{bmatrix} = \frac{1}{32} \begin{bmatrix} 3 & 0 & -3 \\ 10 & 0 & -10 \\ 3 & 0 & -3 \end{bmatrix}$$

$$Scharr_y = \frac{1}{16} \begin{bmatrix} 3 & 10 & 3 \end{bmatrix} * \frac{1}{2} \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} = \frac{1}{32} \begin{bmatrix} 3 & 10 & 3 \\ 0 & 0 & 0 \\ -3 & -10 & -3 \end{bmatrix}$$

Creating larger derivative kernels

If we want a 5x5 Sobel kernel:

1	2	1	*	-1	0	1
---	---	---	---	----	---	---

```
sp.convolve(np.array([1,2,1]),  
            np.array([-1,0,1]), mode='full')
```

-1	-2	0	2	1
----	----	---	---	---

1	2	1	*	1	2	1
---	---	---	---	---	---	---

```
sp.convolve(np.array([1,2,1]),  
            np.array([1,2,1]), mode='full')
```

1	4	6	4	1
---	---	---	---	---

*

`np.outer()` or
`sp.convolve2d()`

-1				

0	0	1	0	0			
0	0	2	0	0			
0	0	0	0	0	0	0	
0	0	-2	0	0	0	0	0
0	0	-1	1	0	4	0	6
				0	0	0	0
				0	0	0	0

We can **generate kernels recursively**, using the 1D central difference kernel and the binomial!

Creating larger derivative kernels

If we want a 5x5 Sobel kernel:

1	2	1
---	---	---

 *

-1	0	1
----	---	---

```
sp.convolve(np.array([1,2,1]),  
            np.array([-1,0,1]), mode='full')
```

-1	-2	0	2	1
----	----	---	---	---

1	2	1
---	---	---

 *

1	2	1
---	---	---

```
sp.convolve(np.array([1,2,1]),  
            np.array([1,2,1]), mode='full')
```

1	4	6	4	1
---	---	---	---	---

*

`np.outer()` or
`sp.convolve2d()`

-1	-4			

0	0	1	0	0		
0	0	2	0	0		
0	0	0	0	0	0	0
0	0	-2	0	0	0	0
0	0	1	-1	4	0	6
		0	0	0	0	0
		0	0	0	0	0

We can **generate kernels recursively**, using the 1D central difference kernel and the binomial!

Creating larger derivative kernels

If we want a 5x5 Sobel kernel:

1	2	1
---	---	---

 *

-1	0	1
----	---	---

```
sp.convolve(np.array([1,2,1]),  
            np.array([-1,0,1]), mode='full')
```

-1	-2	0	2	1
----	----	---	---	---

1	2	1
---	---	---

 *

1	2	1
---	---	---

```
sp.convolve(np.array([1,2,1]),  
            np.array([1,2,1]), mode='full')
```

1	4	6	4	1
---	---	---	---	---

*

`np.outer()` or
`sp.convolve2d()`

-1	-4	-6		

0	0	1	0	0
0	0	2	0	0
0 0	0 0	0 0	0 0	0 0
0 0	0 0	-2 0	0 0	0 0
0 1	0 4	-1 6	0 4	0 1
0	0	0	0	0
0	0	0	0	0

We can **generate kernels recursively**, using the 1D central difference kernel and the binomial!

Creating larger derivative kernels

If we want a 5x5 Sobel kernel:

1	2	1
---	---	---

 *

-1	0	1
----	---	---

```
sp.convolve(np.array([1,2,1]),  
np.array([-1,0,1]), mode='full')
```

-1	-2	0	2	1
----	----	---	---	---

1	2	1
---	---	---

 *

1	2	1
---	---	---

```
sp.convolve(np.array([1,2,1]),  
np.array([1,2,1]), mode='full')
```

1	4	6	4	1
---	---	---	---	---

*

`np.outer()` or
`sp.convolve2d()`

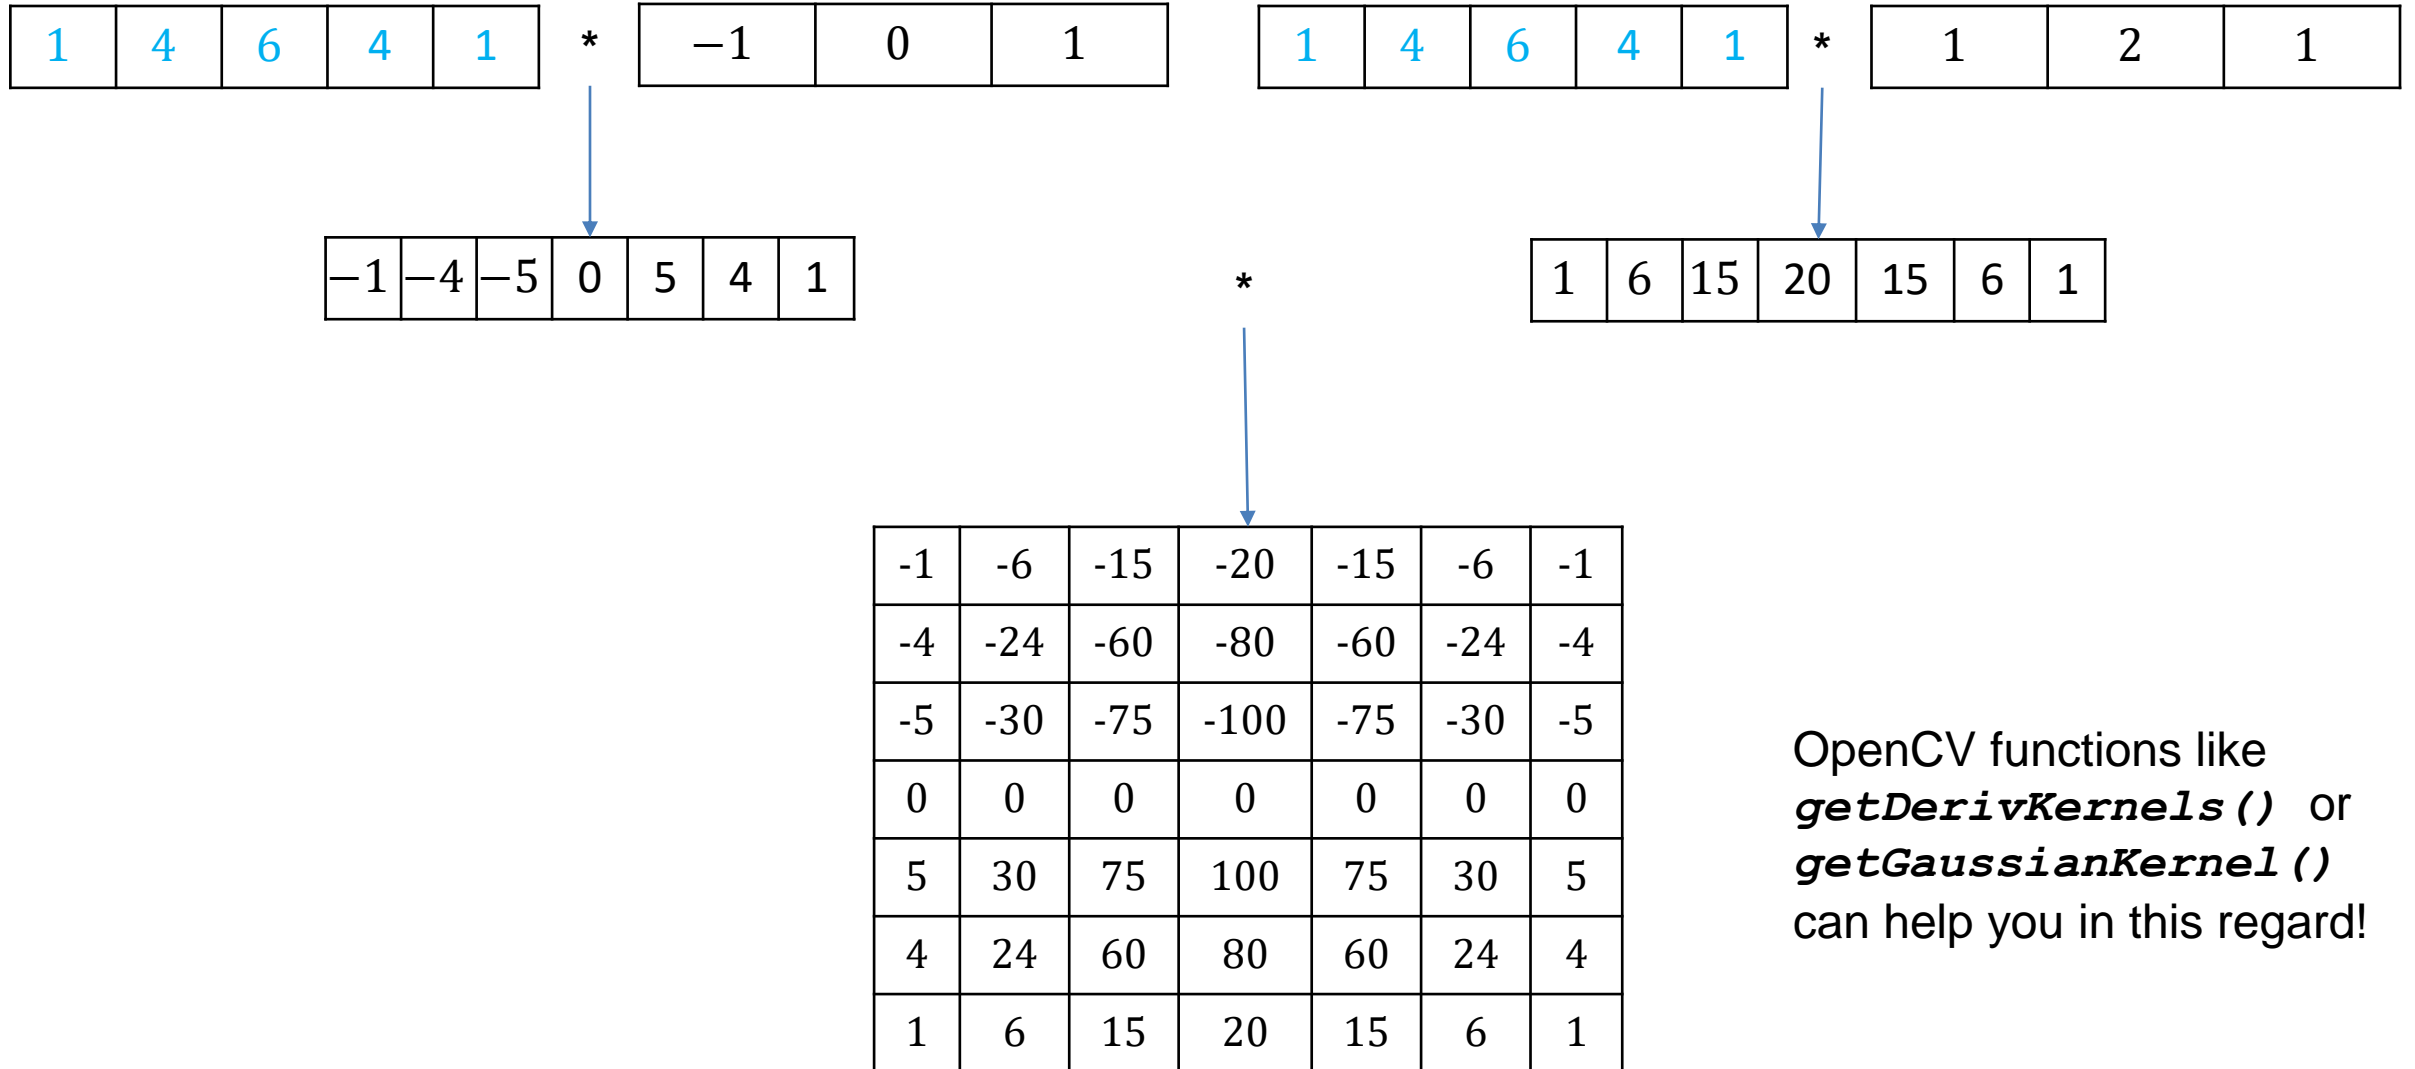
-1	-4	-6	-4	-1
-2	-8	-12	-8	-2
0	0	0	0	0
2	8	12	8	2
1	4	6	4	1

We can **generate kernels recursively**, using the 1D central difference kernel and the binomial!

0	0	0	0	0	
0	0	0	0	0	
1	4	6	4	1	0
0	0	0	0	2	0
0	0	0	0	0	0
		0	0	-2	0
		0	0	-1	0

Creating larger derivative kernels

If we want a 7x7 Sobel kernel:



OpenCV functions like ***getDerivKernels()*** or ***getGaussianKernel()*** can help you in this regard!

Separability

First derivative of 2D Gaussian is separable:

$$\begin{aligned} G(x, y) &= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{(x^2 + y^2)}{2\sigma^2}\right) \\ &= \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{x^2}{2\sigma^2}\right) \cdot \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{y^2}{2\sigma^2}\right) \\ &= G_h(x) \cdot G_v(y) \\ \frac{\partial G(x, y)}{\partial x} &= -\frac{x}{\sigma^2} \cdot \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right) \exp\left(-\frac{y^2}{2\sigma^2}\right) \\ &= G'_h(x) \cdot G_v(y) \end{aligned}$$

horizontal 1D Gaussian derivative kernel vertical 1D Gaussian kernel


smooth in one direction, differentiate in the other

This allows us to know **which 1D kernels to apply, and in what order, to calculate the derivatives of an image.** In the previous example, the first derivative on X.


Laplacian of Gaussian

- How many 2nd derivative filters do we have? There are four 2nd partial derivative filters.
- In practice, it's handy to define a single 2nd derivative filter—the Laplacian
 - Instead of using $\frac{\partial^2(G_\sigma(x,y))}{\partial x^2} = \left(\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2}\right) G_\sigma(x,y)$ and $\frac{\partial^2(G_\sigma(x,y))}{\partial y^2} = \left(\frac{y^2}{\sigma^4} - \frac{1}{\sigma^2}\right) G_\sigma(x,y)$
 - We employ the Laplacian of Gaussian (LoG) filter


Gaussian



1st derivative of Gaussian (x)




LoG




1D Laplacian of Gaussian

Consider $\frac{\partial^2}{\partial x^2}(h \star f)$


f



$\frac{\partial^2}{\partial x^2}h$



$(\frac{\partial^2}{\partial x^2}h) \star f$



Finding these zero-crossings is the goal of the Marr-Hildreth algorithm for edge detection.

Marr, D., & Hildreth, E. (1980). Theory of edge detection. *Procs. of the Royal Society of London*, 207, 187-217.

Where is the edge?


Zero-crossings of bottom graph

Slide credit: Steve Seitz

Laplacian of Gaussian

- Example of 3x3 Laplacian kernel:


0	1	0
1	-4	1
0	1	0



- This kernel is not separable, but...
 - we can calculate LoG as the sum of 2D convolutions that are separable
 - Because the Gaussian and its derivatives are!!!
 - In practice, we need 4 1D convolutions (and that is still less computationally expensive than doing a single 2D convolution; unless the kernel is very small)

if the kernel is 7×7 we need 49 multiplications and additions per pixel for the 2D kernel, or 4·7=28 multiplications and additions per pixel for the four 1D kernels; this difference grows as the kernel gets larger

“inverted Mexican hat”



Laplacian of Gaussian

The **Laplacian operator** is the divergence of the gradient, and is given by the sum of the second order derivatives (i.e. the trace of the Hessian matrix):

$$\boxed{\nabla^2} I = \nabla \cdot \nabla I = \begin{bmatrix} \frac{\partial}{\partial x} & \frac{\partial}{\partial y} \end{bmatrix} \begin{bmatrix} \frac{\partial I}{\partial x} & \frac{\partial I}{\partial y} \end{bmatrix}^T = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

Laplacian of Gaussian

Associative property

Distributive property

$$\underbrace{\frac{\partial^2 (I \circledast G)}{\partial x^2} + \frac{\partial^2 (I \circledast G)}{\partial y^2}}_{\text{Laplacian of an image smoothed with a Gaussian kernel}} = \underbrace{I \circledast \left(\frac{\partial^2 G}{\partial x^2} + \frac{\partial^2 G}{\partial y^2} \right)}_{\text{Image convolved with the Laplacian of a Gaussian}} = I \circledast \frac{\partial^2 G}{\partial x^2} + I \circledast \frac{\partial^2 G}{\partial y^2}$$

Laplacian of Gaussian

How do we compute $I \circledast \frac{\partial^2 G}{\partial x^2} + I \circledast \frac{\partial^2 G}{\partial y^2}$?

We know that the 1st derivative of a 2D Gaussian is separable:

$$\begin{aligned}\frac{\partial G(x, y)}{\partial x} &= -\frac{x}{\sigma^2} \cdot \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right) \exp\left(-\frac{y^2}{2\sigma^2}\right) \\ &= G'_h(x) \cdot G_v(y)\end{aligned}$$

Then:

$$\frac{\partial^2 G}{\partial x^2} + \frac{\partial^2 G}{\partial y^2} = \boxed{G''_h(x) \cdot G_v(y)} + \boxed{G_h(x) \cdot G''_v(y)}$$

convolution in one direction with the 2nd derivative of the Gaussian, and then in the other direction, convolution with the Gaussian.

convolution in one direction with the Gaussian and then in the other, convolution with the 2nd derivative of the Gaussian.






Laplacian of Gaussian

$$L = \sigma^2 \left(G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right) \quad \frac{\partial^2 G}{\partial x^2} + \frac{\partial^2 G}{\partial y^2} = G''_h(x) \cdot G_v(y) + G_h(x) \cdot G''_v(y)$$






Why do we multiply by σ^2 ? 🤔

Intuition: the Gaussian is a probability function, so it has area 1. If we squeeze it in the horizontal direction then we must expand it in the vertical direction to preserve area.

To keep response the same (scale-invariant), must multiply Gaussian derivative by σ
Laplacian is the second Gaussian derivative, so it **must be multiplied by σ^2**



Scale
normalized




Scale
NO normalized


Laplacian of Gaussian

- The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases
- To keep response the same (scale-invariant), must multiply Gaussian derivative by σ
- Laplacian is the second Gaussian derivative, so it must be multiplied by σ^2


Original signal




Unnormalized Laplacian response



Scale-normalized Laplacian response



Responses of convolving
the original signal with the
LoG (using different
sigmas)



Laplacian of Gaussian approximation

- Approximating the Laplacian with a difference of Gaussians:


$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma))$$

(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)

Advantage: we don't even need to implement the second derivative.



Laplacian of Gaussian approximation

- LoG and DoG are **band-pass filters**!


$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma))$$

(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)

The Gaussian functions depicted here do not include the normalizing constant.



FT
↔





Image taken from Orive-Miguel et al. (2019). Improving localization of deep inclusions in time-resolved diffuse optical tomography. *Applied Sciences*, 9(24), 5468.


The Gaussian filters involved are “killing” different high-frequencies. Their **subtraction** give us a range/band of (high) frequencies whose pass is allowed!

Note: We'll understand better this slide once we review the Frequency Domain!

Effect of σ on derivatives (scales)



$\sigma = 1$



$\sigma = 3$

The apparent structures differ depending on Gaussian's scale parameter.

Larger values: larger scale edges detected
Smaller values: finer features detected

Summary: Two types of convolution kernels

Smoothing

$$\sum g_i = 1$$

(Smoothing a constant function should not change it)

Example: $(1/4) * [1 \ 2 \ 1]$

Low-pass filter (Gaussian)

Differentiating

$$\sum g_i = 0$$

(Differentiating a constant function should return 0)

Example: $[-1 \ 0 \ 1]$

High-pass filter (derivative of Gaussian)

... also **bandpass** filters
(Laplacian of Gaussian)

Summary: Two types of convolution kernels

Two closely related problems:



blur
(to remove noise)



differentiate
(to highlight details)

Summary: Two types of convolution kernels

Two closely related problems:



blur
(to remove noise)




differentiate
(to highlight details)

Underlying math is the same!

Canny Edge Detector

- Sometimes we need to accurately locate the boundary....




Canny Edge Detector

- Canny, J., “[A Computational Approach To Edge Detection](#)”, IEEE Trans. on Pattern Analysis and Machine Intelligence, 8(6):679–698, 1986 (~43K citations in Google Scholar)
- Still popular because
 - easy to implement
 - small number of intuitive parameters
 - computationally efficient
 - good results


Python:

https://docs.opencv.org/4.8.0/da/d22/tutorial_py_canny.html

Original Image



Sobel Magnitude




Canny Edge Detector



Canny Edge Detector

- 1) Filter image with derivative of Gaussian and compute the gradient (magnitude and orientation) at each pixel



$$\text{theta} = \text{atan2}(g_y, g_x)$$

360



Gradient orientation angle

0

Canny Edge Detector

2) Non-maximum suppression

A full scan of the image is done to remove any unwanted pixels which may not constitute the edge.



- We check if every pixel is **local maximum in its neighborhood along gradient direction**
 - Point A is on the edge. Gradient direction is normal to the edge. Point B and C are in gradient directions.
 - Point A is checked with point B and C to see if it forms a local maximum.
 - If so, it is considered for next stage, otherwise, it is suppressed (put to zero).

For each point on the edge, we keep the highest (maximum) intensity point.
This makes the edges thinner (1 pixel wide).

Before Non-max Suppression




After Non-max Suppression




Canny Edge Detector

3) Linking and thresholding edges (hysteresis)

- We only want strong edges
- 2 thresholds (high, low), 3 cases (R: response):
 - $R > \text{high}$: strong edge!
 - $R < \text{low}$: no edge
 - $R < \text{high}$ but $R > \text{low}$: weak edge
 - **Weak edges are edges iff they connect to strong ones**
 - Look in some neighborhood (usually 8 closest)



No strong pixels around




One strong pixel around




Canny Edge Detector

3) Linking and thresholding edges (hysteresis)

- We only want strong edges
- 2 thresholds (high, low), 3 cases (R: response):
 - $R > \text{high}$: strong edge!
 - $R < \text{low}$: no edge
 - $R < \text{high}$ but $R > \text{low}$: weak edge
 - **Weak edges are edges iff they connect to strong ones**
 - Look in some neighborhood (usually 8 closest)




No strong pixels around



One strong pixel around



Canny Edge Detector




1. Filter image with derivative of Gaussian and find magnitude and orientation of gradient

2. Non-maximum suppression

3. Linking and thresholding (hysteresis):

- Define two thresholds: low and high
- Use the high threshold to start edge curves and the low threshold to continue them

Canny Edge Detector



Parameters to tune:

high threshold
low threshold

σ : width of the Gaussian blur

- The choice of σ depends on desired behavior
 - Large σ detects “large-scale” edges
 - Small σ detects fine edges

Canny Edge Detector



original image



**high threshold
(strong edges)**



**low threshold
(weak edges)**



hysteresis threshold

Gaussian Filtering and Edges

Pablo Mesejo

pmesejo@go.ugr.es

Universidad de Granada

Departamento de Ciencias de la Computación e Inteligencia Artificial



UNIVERSIDAD
DE GRANADA



DaSCI

Instituto Andaluz de Investigación en
Data Science and Computational Intelligence